

SRL Rubric (100 points) (retired this edition, kept on file)

HONR 46400 · Evidence-Driven Research

⚠ Retired for the Fall 2026 edition: kept for a future one

The Student Research Lead role does not run this semester, so nothing in this document is required of you. Every Monday and Wednesday lecture from Week 2 now opens with a ten-minute **Lab Meeting** instead. One of you is that lecture's **reporter**: you take **seven minutes on a decision from your own project and the evidence behind it**, then **three minutes of questions from the room**. You do not teach the lecture's concept, and **your report is not graded**. The instructor leads from minute 10 and owns accuracy, the AI tooling and the clock.

The 25% **Student Research Lead** grade category retired with the role. In its place, **Lecture Notebooks (20%)** collects the notebook you work in class, once a week. Everyone hands one in, not only the reporter, and it is graded on completion alone.

Nothing below has been removed. This document is kept whole for a future edition that reinstates the role. Where a procedure would mislead you today, it carries a 🛑 **Retired procedure** note naming what happens instead. (*Ruling: [_project_docs/DECISIONS.md](#), D74, 2026-08-31.*)

🛑 **This rubric is not applied in Fall 2026.** Your ten minutes as reporter are ungraded, so no live score is taken against the nine rows below. They are kept exactly as written, every row and every band, for a future edition that reinstates the role. What is graded instead is **Lecture Notebooks (20%)**, on completion only.

How your leading is graded. Read this before your first slot so you know what strong leading looks like, then aim your prepared notebook at it. You are graded live, during the session, on the nine rows below. Each lead is scored on its own, and together those scores make up the SRL share of your course grade.

The rubric measures one thing from nine angles: **did you run a Socratic investigation that kept the room thinking, used AI as an arm rather than a brain, and landed on a defended human decision?** It does not measure whether you knew every answer. Leading an investigation well while genuinely unsure can score Exemplary.

The rubric

Criterion	Exemplary	Proficient	Developing	Beginning
1. Conceptual correctness (15)	The concept in play is handled correctly throughout; any student or AI error is caught and corrected; the room leaves with an accurate understanding (13–15)	Concept is essentially correct; a minor imprecision, self-corrected or harmless (10–12)	A conceptual slip goes uncorrected until the instructor steps in; partial understanding leaves the room (6–9)	A conceptual error spreads unchecked; the room leaves with a wrong idea (0–5)
2. Quality of Socratic questions (15)	Questions are open and non-leading and repeatedly hand the thinking back to the room; the room does the reasoning; the puzzle is staged in the lead's own way rather than read off the brief (13–15)	Mostly strong questions, with an occasional leading or yes/no question (10–12)	Questions often point at the answer or are answerable in a word; the lead does most of the reasoning (6–9)	Little real questioning; the session is a presentation, not an investigation (0–5)

Criterion	Exemplary	Proficient	Developing	Beginning
3. Ability to expose assumptions (10)	Names the hidden assumption under at least one claim and makes the room examine it; the probe is the lesson (9–10)	Surfaces an assumption but does not fully work it (7–8)	Gestures at assumptions without naming one clearly (4–6)	Claims pass unexamined; no assumption surfaced (0–3)
4. Productive use of AI (15)	Commit-before-AI is enforced; AI is pointed at a specific job with a prepared prompt; the room sees AI used as an arm, not a brain (13–15)	AI is used purposefully, with a minor lapse such as a vague prompt or a skipped commitment (10–12)	AI is used loosely; the room reads an AI answer without committing first (6–9)	AI does the thinking, or “explain this topic” is read aloud, or the hard cap below triggers (0–5)
5. Ability to interrogate AI output (15)	Actively hunts the three failure patterns (fabrication, wrong method, missing uncertainty); sources are confirmed; AI overconfidence is called out (13–15)	Interrogates the AI but misses one failure pattern present in the output (10–12)	Light or superficial checking; takes the AI answer largely at face value (6–9)	AI output is accepted without interrogation (0–5)
6. Inclusion of classmates (10)	Every person in the room is drawn in and commits at least once; no one hides (9–10)	Most participate; one voice dominates or one person stays silent (7–8)	A couple of people carry the session; the rest stay passive (4–6)	The lead talks at the room; little participation (0–3)
7. Time management (5)	Hits the section checkpoints; each block gets its minutes; handoffs are clean (5)	A minor overrun, recovered (4)	Misses a checkpoint by a wide margin; a block gets squeezed (2–3)	Runs out of time; sections are dropped (0–1)
8. Connection to research decisions (10)	Explicitly ties the skill to the room’s own projects and the current milestone; a real project decision is named (9–10)	Connects to projects in general terms (7–8)	A token connection tacked on at the end (4–6)	No connection to anyone’s research (0–3)
9. Handling of incorrect or uncertain answers (5)	Wrong answers are examined as data; no one is humiliated; the lead’s own uncertainty becomes an investigation (5)	Wrong answers handled respectfully but moved past quickly (4)	A wrong answer is shut down, or a person is left uncomfortable (2–3)	A classmate is humiliated, or the lead hides from uncertainty by faking an answer (0–1)

Total: 100 points (15 + 15 + 10 + 15 + 15 + 10 + 5 + 10 + 5 = 100).

The hard cap

Presenting an AI answer as settled without verifying it caps row 4 (Productive use of AI) at Beginning (0–5), regardless of the rest of that row.

This is the course’s central discipline with teeth. AI may propose; the researcher must verify. If you let an AI answer stand as the final word in front of the room, without the class confirming sources, cross-checking facts, or committing to a human decision, you have modeled the exact behavior the course exists to prevent, and row 4 cannot exceed Beginning. The fix is simple and lives in every session: run the “After running, verify” checklist out loud, then close on a decision the room defends.

How the score comes together

- You are graded **live**, during your session, against these nine rows.

- The **notebook you submit the day before** feeds the score before you even walk in. A strong  **My Lead Plan**, on top of a notebook you actually worked through, shows up as strong Socratic questions (row 2), a real assumption-probe (row 3), a specific AI plan (rows 4–5), a workable timing plan (row 7), and a named project connection (row 8).
 - **Revision does not apply** to a live performance the way it applies to a written milestone. Your next slot is your improvement window. Read your peer feedback ([srl_peer_feedback_form.md](#)) and the instructor's notes, and bring the fix to the next one.
-

*Related files: [srl_handbook.md](#) (the role these rows measure), [srl_ai_integration_guide.md](#) (rows 4 and 5 in depth), and the  **My Lead Plan** cell in your lecture's notebook (the part that earns most of this score in advance).*