

SRL Submission: What to Hand In

HONR 46400 · Evidence-Driven Research

Read this before your first slot. It is the same for every one of your five leads, so nothing here changes from slot to slot except which lecture you are preparing.

What this assignment collects

The lecture notebook for the lecture you lead, filled in. Only the student holding this SRL slot submits here. If this lecture is not your slot, there is nothing for you to hand in and this assignment does not affect your grade. Open the notebook from the course Schedule page, then **File** → **Save a copy in Drive** so the copy is yours. Two things make it complete. First, fill the  **My Lead Plan** cell at the top of your lecture: how you will stage the puzzle, the one thing you are adding that the SRL Lead Brief does not have, your commitment question and how the room will record it, the decision you will close on with the defense question you will put to one person, and your fallback if the AI tool is slow or down. Second, work the rest of the notebook as a learner, so you have sat with every answer the room will reach before you ask anyone else to reach it.

What to hand in

One notebook. Either paste a Colab share link set so that anyone with the link can comment, which is the better option because it lets me write notes directly on the cells you will teach from, or upload the file itself (**File** → **Download** → **Download .ipynb**). Name it with your last name followed by this assignment's title, for example `LASTNAME_02_studio_lec_01_srl.ipynb`. Submit by the deadline even if the notebook is rough. A rough plan I can send notes on beats a polished one that arrives after the review window has closed.

How it is graded

Your leading itself is scored on the nine rows of the SRL rubric, and the notebook you submit here feeds those rows before you walk in. A strong My Lead Plan shows up as strong Socratic questions, a real assumption probe, a specific AI plan, a workable timing plan, and a named connection to your classmates' projects. Nothing you write here is graded right or wrong.

The rubric your session is scored on

You are graded **live**, during your session, on the nine rows below. Each lead is scored on its own, and together those scores make up the SRL share of your course grade. The rubric measures one thing from nine angles: **did you run a Socratic investigation that kept the room thinking, used AI as an arm rather than a brain, and landed on a defended human decision?** It does not measure whether you knew every answer. Leading an investigation well while genuinely unsure can score Exemplary.

Criterion	Exemplary	Proficient	Developing	Beginning
1. Conceptual correctness (15)	The concept in play is handled correctly throughout; any student or AI error is caught and corrected; the room leaves with an accurate understanding (13–15)	Concept is essentially correct; a minor imprecision, self-corrected or harmless (10–12)	A conceptual slip goes uncorrected until the instructor steps in; partial understanding leaves the room (6–9)	A conceptual error spreads unchecked; the room leaves with a wrong idea (0–5)
2. Quality of Socratic questions (15)	Questions are open and non-leading and repeatedly hand the	Mostly strong questions, with an occasional	Questions often point at the answer or are answerable in a word; the	Little real questioning; the session is a

Criterion	Exemplary	Proficient	Developing	Beginning
	thinking back to the room; the room does the reasoning; the puzzle is staged in the lead's own way rather than read off the brief (13–15)	leading or yes/no question (10–12)	lead does most of the reasoning (6–9)	presentation, not an investigation (0–5)
3. Ability to expose assumptions (10)	Names the hidden assumption under at least one claim and makes the room examine it; the probe is the lesson (9–10)	Surfaces an assumption but does not fully work it (7–8)	Gestures at assumptions without naming one clearly (4–6)	Claims pass unexamined; no assumption surfaced (0–3)
4. Productive use of AI (15)	Commit-before-AI is enforced; AI is pointed at a specific job with a prepared prompt; the room sees AI used as an arm, not a brain (13–15)	AI is used purposefully, with a minor lapse such as a vague prompt or a skipped commitment (10–12)	AI is used loosely; the room reads an AI answer without committing first (6–9)	AI does the thinking, or “explain this topic” is read aloud, or the hard cap below triggers (0–5)
5. Ability to interrogate AI output (15)	Actively hunts the three failure patterns (fabrication, wrong method, missing uncertainty); sources are confirmed; AI overconfidence is called out (13–15)	Interrogates the AI but misses one failure pattern present in the output (10–12)	Light or superficial checking; takes the AI answer largely at face value (6–9)	AI output is accepted without interrogation (0–5)
6. Inclusion of classmates (10)	Every person in the room is drawn in and commits at least once; no one hides (9–10)	Most participate; one voice dominates or one person stays silent (7–8)	A couple of people carry the session; the rest stay passive (4–6)	The lead talks at the room; little participation (0–3)
7. Time management (5)	Hits the section checkpoints; each block gets its minutes; handoffs are clean (5)	A minor overrun, recovered (4)	Misses a checkpoint by a wide margin; a block gets squeezed (2–3)	Runs out of time; sections are dropped (0–1)
8. Connection to research decisions (10)	Explicitly ties the skill to the room's own projects and the current milestone; a real project decision is named (9–10)	Connects to projects in general terms (7–8)	A token connection tacked on at the end (4–6)	No connection to anyone's research (0–3)
9. Handling of incorrect or uncertain answers (5)	Wrong answers are examined as data; no one is humiliated; the lead's own uncertainty becomes an investigation (5)	Wrong answers handled respectfully but moved past quickly (4)	A wrong answer is shut down, or a person is left uncomfortable (2–3)	A classmate is humiliated, or the lead hides from uncertainty by faking an answer (0–1)

Total: 100 points (15 + 15 + 10 + 15 + 15 + 10 + 5 + 10 + 5 = 100).

The hard cap

Presenting an AI answer as settled without verifying it caps row 4 (Productive use of AI) at Beginning (0–5), regardless of the rest of that row.

This is the course's central discipline with teeth. AI may propose; the researcher must verify. If you let an AI answer stand as the final word in front of the room, without the class confirming sources, cross-checking facts, or committing to a human

decision, you have modeled the exact behavior the course exists to prevent, and row 4 cannot exceed Beginning. The fix is simple and lives in every session: run the “After running, verify” checklist out loud, then close on a decision the room defends.

The rubric above is reproduced from [srl_rubric.md](#), which is the canonical version and carries the fuller explanation of how the score comes together. Related: [srl_handbook.md](#) (the role these rows measure), [srl_ai_integration_guide.md](#) (rows 4 and 5 in depth), [socratic_question_bank.md](#) (questions to reach for when you get stuck).